

# Machine Learning Homework 4\_Report

2017112167 이수진

## 1. CNN & MNIST

<full program code> : I wrote only the red part for each question.

```
import torch
import torch.nn as nn
from torch.autograd import Variable
import torch.utils.data as Data
import torchvision
import matplotlib.pyplot as plt
#matplotlib inline

torch.manual_seed(1)    # reproducible
# Hyper Parameters
EPOCH = 1                # train the training data n times, to save time, we just train 1
epoch
BATCH_SIZE = 50
LR = 0.001                # learning rate
DOWNLOAD_MNIST = False   # set to False if you have downloaded

# Mnist digits dataset
train_data = torchvision.datasets.MNIST(
    root='./mnist/',
    train=True,
    transform=torchvision.transforms.ToTensor(),    # Converts a PIL.Image or numpy.ndarray t
o
                                                    # torch.FloatTensor of shape (C x
H x W) and normalize in the range [0.0, 1.0]
    download=False,    # download it if you don't have it
)
# plot one example
print(train_data.train_data.size())    # (60000, 28, 28)
print(train_data.train_labels.size())    # (60000)
plt.imshow(train_data.train_data[0].numpy(), cmap='gray')
plt.title('%i' % train_data.train_labels[0])
plt.show()

# Data Loader for easy mini-batch return in training, the image batch shape will be (50, 1,
28, 28)
train_loader = Data.DataLoader(dataset=train_data, batch_size=BATCH_SIZE, shuffle=True)

# convert test data into Variable, pick 2000 samples to speed up testing
test_data = torchvision.datasets.MNIST(root='./mnist/', train=False)
print(test_data.test_data[0].size())
test_x = Variable(torch.unsqueeze(test_data.test_data, dim=1)).type(torch.FloatTensor)[:2000]/255
.
print(test_x[0].size())
# shape from (2000, 28, 28) to (2000, 1, 28, 28), value in range(0,1)
test_y = test_data.test_labels[:2000]
print(test_y[0])

class CNN(nn.Module):
    def __init__(self):
```

```

super(CNN, self).__init__()
self.conv1 = nn.Sequential(          # input shape (1, 28, 28)
    nn.Conv2d(
        in_channels=1,              # input height
        out_channels=16,            # n_filters
        kernel_size=5,              # filter size
        stride=1,                   # filter movement/step
        padding=2,                  # if want same width and length of this im
age after con2d, padding=(kernel_size-1)/2 if stride=1
    ),                               # output shape (16, 28, 28)
    nn.ReLU(),                      # activation
    nn.MaxPool2d(kernel_size=2),    # choose max value in 2x2 area, output shape
(16, 14, 14)
)
self.conv2 = nn.Sequential(          # input shape (1, 28, 28)
    nn.Conv2d(16, 32, 5, 1, 2),    # output shape (32, 14, 14)
    nn.ReLU(),                     # activation
    nn.MaxPool2d(2),               # output shape (32, 7, 7)
)
self.out = nn.Linear(32 * 7 * 7, 10) # fully connected layer, output 10 classes

def forward(self, x):
    x = self.conv1(x)
    x = self.conv2(x)
    x = x.view(x.size(0), -1)       # flatten the output of conv2 to (batch_size, 32
* 7 * 7)
    output = self.out(x)
    return output, x                # return x for visualization

cnn = CNN()
print(cnn) # net architecture

optimizer = torch.optim.Adam(cnn.parameters(), lr=LR) # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()                    # the target label is not one-hotte
d

# following function (plot_with_labels) is for visualization, can be ignored if not interested
from matplotlib import cm
try: from sklearn.manifold import TSNE; HAS_SK = True
except: HAS_SK = False; print('Please install sklearn for layer visualization')
def plot_with_labels(lowDWeights, labels):
    plt.cla()
    X, Y = lowDWeights[:, 0], lowDWeights[:, 1]
    for x, y, s in zip(X, Y, labels):
        c = cm.rainbow(int(255 * s / 9)); plt.text(x, y, s, backgroundcolor=c, fontsize=9)
    plt.xlim(X.min(), X.max()); plt.ylim(Y.min(), Y.max()); plt.title('Visualize last layer'); plt.show();
    plt.pause(0.01)

plt.ion()
# training and testing
for epoch in range(EPOCH):
    for step, (x, y) in enumerate(train_loader): # gives batch data, normalize x when itera
te train_loader
        b_x = Variable(x) # batch x
        b_y = Variable(y) # batch y

        output = cnn(b_x)[0] # cnn output

```

```

loss = loss_func(output, b_y) # cross entropy loss
optimizer.zero_grad()         # clear gradients for this training step
loss.backward()                # backpropagation, compute gradients
optimizer.step()               # apply gradients

if step % 100 == 0:
    test_output, last_layer = cnn(test_x)
    pred_y = torch.max(test_output, 1)[1].data.squeeze()
    accuracy = (pred_y == test_y).sum().item() / float(test_y.size(0))
    print('Epoch: ', epoch, '| train loss: %.4f' % loss.data, '| test accuracy: %.2f' %
accuracy)
    if HAS_SK:
        # Visualization of trained flatten layer (T-SNE)
        tsne = TSNE(perplexity=30, n_components=2, init='pca', n_iter=5000)
        plot_only = 500
        low_dim_embs = tsne.fit_transform(last_layer.data.numpy()[:plot_only, :])
        labels = test_y.numpy()[:plot_only]
        plot_with_labels(low_dim_embs, labels)

plt.ioff()

```

2) change the current kernel size of the program to different size. (Change 'kernel\_size' parameter of 'Conv2D' function.) Repeat this three times and compare the results.

#### #1-2 kernel\_size=7

```

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(
                in_channels=1,                # input height
                out_channels=16,              # n_filters
                kernel_size=7,                # filter size
                stride=1,                     # filter movement/step
                padding=3,                    # if want same width and length of this image
                ),                             # output shape (16, 28, 28)
            nn.ReLU(),                       # activation
            nn.MaxPool2d(kernel_size=2),      # choose max value in 2x2 area, output shape (16,
14, 14)
        )
        self.conv2 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(16, 32, 7, 1, 3),      # output shape (32, 14, 14)
            nn.ReLU(),                       # activation
            nn.MaxPool2d(2),                 # output shape (32, 7, 7)
        )
        self.out = nn.Linear(32 * 7 * 7, 10) # fully connected layer, output 10 classes


    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = x.view(x.size(0), -1)           # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
        output = self.out(x)
        return output, x                    # return x for visualization

cnn = CNN()
print(cnn) # net architecture

```

```
optimizer = torch.optim.Adam(cnn.parameters(), lr=LR)    # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()                       # the target label is not one-hot
```

Epoch: 0 | train loss: 0.0769 | test accuracy: 0.98



## #1-2 kernel\_size= 5

```
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(                    # input shape (1, 28, 28)
            nn.Conv2d(                                # input height
                in_channels=1,                          # n_filters
                out_channels=16,                         # filter size
                kernel_size=5,                           # filter movement/step
                stride=1,                                # if want same width and length of this image
                padding=2,
            ),                                           # output shape (16, 28, 28)
            nn.ReLU(),                                # activation
            nn.MaxPool2d(kernel_size=2),                # choose max value in 2x2 area, output shape (16,
14, 14)
        )
        self.conv2 = nn.Sequential(                    # input shape (1, 28, 28)
            nn.Conv2d(16, 32, 5, 1, 2),                # output shape (32, 14, 14)
            nn.ReLU(),                                # activation
            nn.MaxPool2d(2),                            # output shape (32, 7, 7)
        )
        self.out = nn.Linear(32 * 7 * 7, 10)          # fully connected layer, output 10 classes


    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = x.view(x.size(0), -1)                    # flatten the output of conv2 to (batch_size, 32 * 7 *
7)
        output = self.out(x)
        return output, x    # return x for visualization

cnn = CNN()
print(cnn)    # net architecture

optimizer = torch.optim.Adam(cnn.parameters(), lr=LR)    # optimize all cnn parameters
```

```
loss_func = nn.CrossEntropyLoss() # the target label is not one-hot
```

```
Epoch: 0 | train loss: 0.0257 | test accuracy: 0.98
```



### #1-2 kernel\_size= 3


```
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential( # input shape (1, 28, 28)
            nn.Conv2d( # input height
                in_channels=1, # n_filters
                out_channels=16, # filter size
                kernel_size=3, # filter movement/step
                stride=1, # if want same width and length of this image
                padding=1, # output shape (16, 28, 28)
            ), # activation
            nn.ReLU(), # choose max value in 2x2 area, output shape (16,
            nn.MaxPool2d(kernel_size=2), 14, 14)
        )
        self.conv2 = nn.Sequential( # input shape (1, 28, 28)
            nn.Conv2d(16, 32, 3, 1, 1), # output shape (32, 14, 14)
            nn.ReLU(), # activation
            nn.MaxPool2d(2), # output shape (32, 7, 7)
        )
        self.out = nn.Linear(32 * 7 * 7, 10) # fully connected layer, output 10 classes

    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = x.view(x.size(0), -1) # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
        output = self.out(x)
        return output, x # return x for visualization

cnn = CNN()
print(cnn) # net architecture

optimizer = torch.optim.Adam(cnn.parameters(), lr=LR) # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss() # the target label is not one-hot
```

Epoch: 0 | train loss: 0.1776 | test accuracy: 0.97



---

### Report

- ⇒ When kernel size is 7, 5, or 5, the comparison was made. Among the results with the highest test accuracy, the graph with the lowest train loss was selected and compared. Accuracy increased as kernel size increased, and train loss increased as kernel size decreased.
- 

3) remove pooling layer in the program (you can remove 'MaxPool2D' function) and compare the results


```
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(
                in_channels=1,
                out_channels=16,
                kernel_size=5,
                stride=1,
                padding=2,
            ),
            nn.ReLU(),
        )
        self.conv2 = nn.Sequential(
            nn.Conv2d(16, 32, 5, 1, 2),
            nn.ReLU(),
        )
        self.out = nn.Linear(32 * 28 * 28, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = x.view(x.size(0), -1)
        output = self.out(x)
        return output, x    # return x for visualization

cnn = CNN()
print(cnn)    # net architecture

optimizer = torch.optim.Adam(cnn.parameters(), lr=LR)    # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()    # the target label is not one-hot
```

Epoch: 0 | train loss: 0.0493 | test accuracy: 0.98



### Report

⇒ The accuracy is similar, but the train loss increases.

4) change the current activation function to other activation function (e.g. sigmoid, tanh, etc). You can do so by `nn.Sigmoid()` to `nn.ReLU()`, `nn.Tanh()`, etc) Repeat this three times and compare the results.

#### #1-4 ReLU()

```
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(                        # input height
                1,                            # n_filters
                16,                           # filter size
                5,                             # filter movement/step
                1,                             # if want same width and length of this image
                padding=2,                    # after conv2d, padding=(kernel_size-1)/2 if stride=1
            ),                                # output shape (16, 28, 28)
            nn.ReLU(),                        # activation
            nn.MaxPool2d(kernel_size=2),      # choose max value in 2x2 area, output shape (16,
14, 14)
        )
        self.conv2 = nn.Sequential(         # input shape (1, 28, 28)
            nn.Conv2d(16, 32, 5, 1, 2),     # output shape (32, 14, 14)
            nn.ReLU(),                      # activation
            nn.MaxPool2d(2),                # output shape (32, 7, 7)
        )
        self.out = nn.Linear(32 * 7 * 7, 10) # fully connected layer, output 10 classes

    def forward(self, x):
        x = self.conv1(x)
```


```

        x = self.conv2(x)
        x = x.view(x.size(0), -1)          # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
        output = self.out(x)
        return output, x                    # return x for visualization
    cnn = CNN()
    print(cnn)    # net architecture

    optimizer = torch.optim.Adam(cnn.parameters(), lr=LR)    # optimize all cnn parameters
    loss_func = nn.CrossEntropyLoss()                        # the target label is not one-hot

```

Epoch: 0 | train loss: 0.0257 | test accuracy: 0.98



#### #1-4 Tanh()

```

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(                        # input height
                in_channels=1,                # n_filters
                out_channels=16,               # filter size
                kernel_size=5,                # filter movement/step
                stride=1,                     # if want same width and length of this image
                padding=2,
                after con2d, padding=(kernel_size-1)/2 if stride=1
            ),                                # output shape (16, 28, 28)
            nn.Tanh(),                        # activation
            nn.MaxPool2d(kernel_size=2),      # choose max value in 2x2 area, output shape (16,
14, 14)
        )
        self.conv2 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(16, 32, 5, 1, 2),      # output shape (32, 14, 14)
            nn.Tanh(),                        # activation
            nn.MaxPool2d(2),                 # output shape (32, 7, 7)
        )
        self.out = nn.Linear(32 * 7 * 7, 10) # fully connected layer, output 10 classes

```



```


def forward(self, x):
    x = self.conv1(x)
    x = self.conv2(x)
    x = x.view(x.size(0), -1)          # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
    output = self.out(x)
    return output, x    # return x for visualization

cnn = CNN()
print(cnn)    # net architecture

optimizer = torch.optim.Adam(cnn.parameters(), lr=LR)    # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()    # the target label is not one-hot

Epoch:  0 | train loss: 0.0224 | test accuracy: 0.98

```



#### #1-4 Sigmoid()

```

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(
                in_channels=1,                # input height
                out_channels=16,              # n_filters
                kernel_size=5,                # filter size
                stride=1,                     # filter movement/step
                padding=2,                    # if want same width and length of this image
            ),                                # output shape (16, 28, 28)
            nn.Sigmoid(),                    # activation
            nn.MaxPool2d(kernel_size=2),      # choose max value in 2x2 area, output shape (16,
14, 14)
        )
        self.conv2 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(16, 32, 5, 1, 2),      # output shape (32, 14, 14)

```

```

        nn.Sigmoid(),                # activation
        nn.MaxPool2d(2),            # output shape (32, 7, 7)
    )
    self.out = nn.Linear(32 * 7 * 7, 10) # fully connected layer, output 10 classes


    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = x.view(x.size(0), -1)      # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
        output = self.out(x)
        return output, x               # return x for visualization

cnn = CNN()
print(cnn) # net architecture

optimizer = torch.optim.Adam(cnn.parameters(), lr=LR) # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()                     # the target label is not one-hot

```

Epoch: 0 | train loss: 0.1897 | test accuracy: 0.94



## Report

⇒ When using Tanh() and ReLU(), the accuracy is higher and the train loss is smaller than when using Sigmoid(). Sigmoid() has poor function.

5) change the current optimization method to other optimization methods (e.g. adam, adaGrad, RMSProp, adaDelta, etc). You can use torch.optim.Adam, etc. Repeat this three times and compare the results.

## #1-5 Adagrad

```

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(
                in_channels=1,                # input height
                out_channels=16,              # n_filters
                kernel_size=5,                # filter size
                stride=1,                     # filter movement/step
                padding=2,                    # if want same width and length of this image after conv2d,

```

```

padding=(kernel_size-1)/2 if stride=1
    ),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2),
)
self.conv2 = nn.Sequential(
    nn.Conv2d(16, 32, 5, 1, 2),
    nn.ReLU(),
    nn.MaxPool2d(2),
)
self.out = nn.Linear(32 * 7 * 7, 10)

def forward(self, x):
    x = self.conv1(x)
    x = self.conv2(x)
    x = x.view(x.size(0), -1)
    output = self.out(x)
    return output, x


cnn = CNN()
print(cnn)

optimizer = torch.optim.Adagrad(cnn.parameters(), lr=LR)
loss_func = nn.CrossEntropyLoss()

```

# output shape (16, 28, 28)  
# activation  
# choose max value in 2x2 area, output shape (16, 14, 14)  
# input shape (1, 28, 28)  
# output shape (32, 14, 14)  
# activation  
# output shape (32, 7, 7)  
# fully connected layer, output 10 classes  
# flatten the output of conv2 to (batch\_size, 32 \* 7 \* 7)  
# return x for visualization  
# net architecture  
# optimize all cnn parameters  
# the target label is not one-hot

Epoch: 0 | train loss: 0.2532 | test accuracy: 0.88



## #1-5 RMSprop

```

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(
                in_channels=1,
                out_channels=16,
                kernel_size=5,
                stride=1,
                padding=2,
            ),
        )

```

# input shape (1, 28, 28)  
# input height  
# n\_filters  
# filter size  
# filter movement/step  
# if want same width and length of this image after con2d,  
# output shape (16, 28, 28)

```

        nn.ReLU(),                # activation
        nn.MaxPool2d(kernel_size=2), # choose max value in 2x2 area, output shape (16, 14, 14)
    )
    self.conv2 = nn.Sequential(    # input shape (1, 28, 28)
        nn.Conv2d(16, 32, 5, 1, 2), # output shape (32, 14, 14)
        nn.ReLU(),                # activation
        nn.MaxPool2d(2),          # output shape (32, 7, 7)
    )
    self.out = nn.Linear(32 * 7 * 7, 10) # fully connected layer, output 10 classes


def forward(self, x):
    x = self.conv1(x)
    x = self.conv2(x)
    x = x.view(x.size(0), -1)      # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
    output = self.out(x)
    return output, x              # return x for visualization

cnn = CNN()
print(cnn) # net architecture

optimizer = torch.optim.RMSprop(cnn.parameters(), lr=LR) # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()

```

Epoch: 0 | train loss: 0.0280 | test accuracy: 0.98



## #1-5 Adadelta

```

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(    # input shape (1, 28, 28)
            nn.Conv2d(
                in_channels=1,          # input height
                out_channels=16,         # n_filters
                kernel_size=5,           # filter size
                stride=1,                # filter movement/step
                padding=2,               # if want same width and length of this image after con2d,
            ),                          # output shape (16, 28, 28)
            nn.ReLU(),                 # activation
            nn.MaxPool2d(kernel_size=2), # choose max value in 2x2 area, output shape (16, 14, 14)
        )

```

```

self.conv2 = nn.Sequential(          # input shape (1, 28, 28)
    nn.Conv2d(16, 32, 5, 1, 2),      # output shape (32, 14, 14)
    nn.ReLU(),                      # activation
    nn.MaxPool2d(2),                 # output shape (32, 7, 7)
)
self.out = nn.Linear(32 * 7 * 7, 10) # fully connected layer, output 10 classes


def forward(self, x):
    x = self.conv1(x)
    x = self.conv2(x)
    x = x.view(x.size(0), -1)        # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
    output = self.out(x)
    return output, x                 # return x for visualization

cnn = CNN()
print(cnn) # net architecture

optimizer = torch.optim.Adadelta(cnn.parameters(), lr=LR) # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()

```

Epoch: 0 | train loss: 2.2464 | test accuracy: 0.34



## Report

⇒ Adadelta's accuracy is very poor. Adagrad is more accurate than Adadelta, but not as good. RMSprop is much more accurate than others and the train loss is very small.

6) now add the Xavier weight initialization method and compare the results. (use torch.nn.init.xavier\_uniform)

```

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(
                in_channels=1,                # input height
                out_channels=16,              # n_filters
                kernel_size=5,                # filter size
                stride=1,                     # filter movement/step
                padding=2,                    # if want same width and length of this image after conv2d,
                padding=(kernel_size-1)/2 if stride=1
            ),                                # output shape (16, 28, 28)
            nn.ReLU(),                       # activation

```

```

        nn.MaxPool2d(kernel_size=2),      # choose max value in 2x2 area, output shape (16, 14, 14)
    )
    self.conv2 = nn.Sequential(           # input shape (1, 28, 28)
        nn.Conv2d(16, 32, 5, 1, 2),      # output shape (32, 14, 14)
        nn.ReLU(),                       # activation
        nn.MaxPool2d(2),                 # output shape (32, 7, 7)
    )
    self.out = nn.Linear(32 * 7 * 7, 10)  # fully connected layer, output 10 classes
    torch.nn.init.xavier_uniform_(self.out.weight)

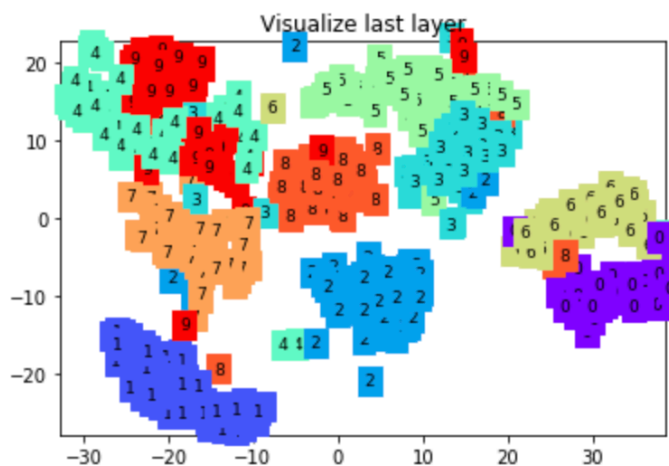
    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = x.view(x.size(0), -1)        # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
        output = self.out(x)
        return output, x                 # return x for visualization

cnn = CNN()
print(cnn)    # net architecture

optimizer = torch.optim.Adam(cnn.parameters(), lr=LR)    # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()                        # the target label is not one-hot

```

Epoch: 0 | train loss: 0.1222 | test accuracy: 0.98



### Report

⇒ The accuracy is similar, but the train loss increases.

7) choose ONE other parameters of CNN program (e.g. number of hidden nodes, dropout, epochs, batch normalization, etc). Change the value of this parameter and compare the results.

### #1-7 EPOCH = 2

```

import torch
import torch.nn as nn
from torch.autograd import Variable
import torch.utils.data as Data
import torchvision
import matplotlib.pyplot as plt
#matplotlib inline

```

```

torch.manual_seed(1)    # reproducible
# Hyper Parameters
EPOCH = 2                # train the training data n times, to save time, we just train 1 epoch
BATCH_SIZE = 50
LR = 0.001                # learning rate
DOWNLOAD_MNIST = False   # set to False if you have downloaded

# Mnist digits dataset
train_data = torchvision.datasets.MNIST(
    root='./mnist/',
    train=True,                # this is training data
    transform=torchvision.transforms.ToTensor(),    # Converts a PIL.Image or numpy.ndarray to
                                                # torch.FloatTensor of shape (C x H x W) and
normalize in the range [0.0, 1.0]
    download=False,            # download it if you don't have it
)
# plot one example
print(train_data.train_data.size())        # (60000, 28, 28)
print(train_data.train_labels.size())      # (60000)
plt.imshow(train_data.train_data[0].numpy(), cmap='gray')
plt.title('%i' % train_data.train_labels[0])
plt.show()

# Data Loader for easy mini-batch return in training, the image batch shape will be (50, 1, 28, 28)
train_loader = Data.DataLoader(dataset=train_data, batch_size=BATCH_SIZE, shuffle=True)

# convert test data into Variable, pick 2000 samples to speed up testing
test_data = torchvision.datasets.MNIST(root='./mnist/', train=False)
print(test_data.test_data[0].size())
test_x = Variable(torch.unsqueeze(test_data.test_data, dim=1)).type(torch.FloatTensor)[:2000]/255.
print(test_x[0].size())
# shape from (2000, 28, 28) to (2000, 1, 28, 28), value in range(0,1)
test_y = test_data.test_labels[:2000]
print(test_y[0])

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Sequential(            # input shape (1, 28, 28)
            nn.Conv2d(
                in_channels=1,                # input height
                out_channels=16,              # n_filters
                kernel_size=5,                # filter size
                stride=1,                     # filter movement/step
                padding=2,                    # if want same width and length of this image after conv2d,
            ),                                # output shape (16, 28, 28)
            nn.ReLU(),                        # activation
            nn.MaxPool2d(kernel_size=2),      # choose max value in 2x2 area, output shape (16, 14, 14)
        )
        self.conv2 = nn.Sequential(          # input shape (1, 28, 28)
            nn.Conv2d(16, 32, 5, 1, 2),      # output shape (32, 14, 14)
            nn.ReLU(),                       # activation
            nn.MaxPool2d(2),                 # output shape (32, 7, 7)
        )
        self.out = nn.Linear(32 * 7 * 7, 10) # fully connected layer, output 10 classes

```

```

def forward(self, x):
    x = self.conv1(x)
    x = self.conv2(x)
    x = x.view(x.size(0), -1)          # flatten the output of conv2 to (batch_size, 32 * 7 * 7)
    output = self.out(x)
    return output, x          # return x for visualization

cnn = CNN()
print(cnn)  # net architecture

optimizer = torch.optim.Adam(cnn.parameters(), lr=LR)  # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()                    # the target label is not one-hot

# following function (plot_with_labels) is for visualization, can be ignored if not interested
from matplotlib import cm
try: from sklearn.manifold import TSNE; HAS_SK = True
except: HAS_SK = False; print('Please install sklearn for layer visualization')
def plot_with_labels(lowDWeights, labels):
    plt.cla()
    X, Y = lowDWeights[:, 0], lowDWeights[:, 1]
    for x, y, s in zip(X, Y, labels):
        c = cm.rainbow(int(255 * s / 9)); plt.text(x, y, s, backgroundcolor=c, fontsize=9)
    plt.xlim(X.min(), X.max()); plt.ylim(Y.min(), Y.max()); plt.title('Visualize last layer'); plt.show(); plt.pause(0.01)

plt.ion()
# training and testing
for epoch in range(EPOCH):
    for step, (x, y) in enumerate(train_loader):  # gives batch data, normalize x when iterate train_loader
        b_x = Variable(x)  # batch x
        b_y = Variable(y)  # batch y

        output = cnn(b_x)[0]  # cnn output
        loss = loss_func(output, b_y)  # cross entropy loss
        optimizer.zero_grad()  # clear gradients for this training step
        loss.backward()  # backpropagation, compute gradients
        optimizer.step()  # apply gradients


    if step % 100 == 0:
        test_output, last_layer = cnn(test_x)
        pred_y = torch.max(test_output, 1)[1].data.squeeze()
        accuracy = (pred_y == test_y).sum().item() / float(test_y.size(0))
        print('Epoch: ', epoch, '| train loss: %.4f' % loss.data, '| test accuracy: %.2f' % accuracy)
        if HAS_SK:
            # Visualization of trained flatten layer (T-SNE)
            tsne = TSNE(perplexity=30, n_components=2, init='pca', n_iter=5000)
            plot_only = 500
            low_dim_embs = tsne.fit_transform(last_layer.data.numpy()[plot_only, :])
            labels = test_y.numpy()[plot_only]
            plot_with_labels(low_dim_embs, labels)

plt.ioff()

```



Epoch: 1 | train loss: 0.0023 | test accuracy: 0.99



### Report

⇒ I changed EPOCH as 2. The test accuracy is increased, and train loss decreased. This is the best result.

## 2. CNN & CIFAR-10

<full program code> : I wrote only the red part for each question.

```
import torch
import torchvision
import torchvision.transforms as transforms

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

import matplotlib.pyplot as plt
import numpy as np

# functions to show an image
def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))

import torch.nn as nn
import torch.nn.functional as F
```

```

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 16, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(16, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
net.to(device)

import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

print('start traning.')
for epoch in range(2):  # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999:  # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
            running_loss = 0.0

```

```

print('Finished Training')

correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (
    100 * correct / total))

```

```

class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
        for i in range(4):
            label = labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1

for i in range(10):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))

```

# del dataiter

2) kernal size

```

#2-2 kernal size = 5
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 16, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(16, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)

```

```
return x
```

```
net = Net()
```

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")  
print(device)  
net.to(device)
```

```
import torch.optim as optim
```

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```
cpu  
start traning.  
[1, 2000] loss: 2.165  
[1, 4000] loss: 1.801  
[1, 6000] loss: 1.598  
[1, 8000] loss: 1.514  
[1, 10000] loss: 1.429  
[1, 12000] loss: 1.389  
[2, 2000] loss: 1.324  
[2, 4000] loss: 1.295  
[2, 6000] loss: 1.270  
[2, 8000] loss: 1.234  
[2, 10000] loss: 1.214  
[2, 12000] loss: 1.183  
Finished Training  
Accuracy of the network on the 10000 test images: 59 %  
Accuracy of plane : 60 %  
Accuracy of car : 79 %  
Accuracy of bird : 36 %  
Accuracy of cat : 26 %  
Accuracy of deer : 62 %  
Accuracy of dog : 55 %  
Accuracy of frog : 69 %  
Accuracy of horse : 70 %  
Accuracy of ship : 72 %  
Accuracy of truck : 57 %
```

```
#2-2 kernal size = 7
```

```
class Net(nn.Module):
```

```
    def __init__(self):  
        super(Net, self).__init__()  
        self.conv1 = nn.Conv2d(3, 16, 7)  
        self.pool = nn.MaxPool2d(2, 2)  
        self.conv2 = nn.Conv2d(16, 16, 7)  
        self.fc1 = nn.Linear(16 * 3 * 3, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)
```

```
    def forward(self, x):  
        x = self.pool(F.relu(self.conv1(x)))  
        x = self.pool(F.relu(self.conv2(x)))  
        x = x.view(-1, 16 * 3 * 3)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)
```

```
return x
```

```
net = Net()
```

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")  
print(device)  
net.to(device)
```

```
import torch.optim as optim
```

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```
cpu  
start traning.  
[1, 2000] loss: 2.182  
[1, 4000] loss: 1.817  
[1, 6000] loss: 1.688  
[1, 8000] loss: 1.588  
[1, 10000] loss: 1.517  
[1, 12000] loss: 1.486  
[2, 2000] loss: 1.413  
[2, 4000] loss: 1.419  
[2, 6000] loss: 1.349  
[2, 8000] loss: 1.352  
[2, 10000] loss: 1.346  
[2, 12000] loss: 1.292  
Finished Training  
Accuracy of the network on the 10000 test images: 54 %  
Accuracy of plane : 51 %  
Accuracy of car : 72 %  
Accuracy of bird : 44 %  
Accuracy of cat : 31 %  
Accuracy of deer : 33 %  
Accuracy of dog : 42 %  
Accuracy of frog : 73 %  
Accuracy of horse : 64 %  
Accuracy of ship : 62 %  
Accuracy of truck : 65 %
```

## Report

> Kernal size 7 is better accuracy than 5

## 3. RNN & Mnist

<full program code> : I wrote only the red part for each question.

```
import torch  
from torch import nn  
from torch.autograd import Variable  
import torchvision.datasets as dsets  
import torchvision.transforms as transforms  
import matplotlib.pyplot as plt
```

```
torch.manual_seed(1)    # reproducible
```

```
# Hyper Parameters
```

```
EPOCH = 1                # train the training data n times, to save time, we just train 1 epoch
```

```
BATCH_SIZE = 64
```

```
TIME_STEP = 28           # rnn time step / image height
```

```

INPUT_SIZE = 28          # rnn input size / image width
LR = 0.01                # learning rate
DOWNLOAD_MNIST = True    # set to True if haven't download the data

# Mnist digital dataset
train_data = datasets.MNIST(
    root='./mnist/',
    train=True,          # this is training data
    transform=transforms.ToTensor(), # Converts a PIL.Image or numpy.ndarray to
                                   # torch.FloatTensor of shape (C x H x W) and normalize in the
    range=[0.0, 1.0]
    download=DOWNLOAD_MNIST, # download it if you don't have it
)

# plot one example
print(train_data.train_data.size()) # (60000, 28, 28)
print(train_data.train_labels.size()) # (60000)
plt.imshow(train_data.train_data[0].numpy(), cmap='gray')
plt.title('%i' % train_data.train_labels[0])
plt.show()

# Data Loader for easy mini-batch return in training
train_loader = torch.utils.data.DataLoader(dataset=train_data, batch_size=BATCH_SIZE, shuffle=True)

# convert test data into Variable, pick 2000 samples to speed up testing
test_data = datasets.MNIST(root='./mnist/', train=False, transform=transforms.ToTensor())
test_x = Variable(test_data.test_data, volatile=True).type(torch.FloatTensor)[:2000]/255. # shape (2000, 28,
28) value in range(0,1)
test_y = test_data.test_labels.numpy().squeeze()[:2000] # covert to numpy array

class RNN(nn.Module):
    def __init__(self):
        super(RNN, self).__init__()

        self.rnn = nn.LSTM( # if use nn.RNN(), it hardly learns
            input_size=INPUT_SIZE,
            hidden_size=64, # rnn hidden unit
            num_layers=1, # number of rnn layer
            batch_first=True, # input & output will has batch size as 1s dimension. e.g. (batch,
time_step, input_size)
        )

        self.out = nn.Linear(64, 10)

    def forward(self, x):
        # x shape (batch, time_step, input_size)
        # r_out shape (batch, time_step, output_size)
        # h_n shape (n_layers, batch, hidden_size)
        # h_c shape (n_layers, batch, hidden_size)
        r_out, (h_n, h_c) = self.rnn(x, None) # None represents zero initial hidden state

        # choose r_out at the last time step
        out = self.out(r_out[:, -1, :])
        return out

rnn = RNN()
print(rnn)

```

```

optimizer = torch.optim.Adam(rnn.parameters(), lr=LR) # optimize all rnn parameters
loss_func = nn.CrossEntropyLoss() # the target label is not one-hot

# training and testing
for epoch in range(EPOCH):
    for step, (x, y) in enumerate(train_loader): # gives batch data
        b_x = Variable(x.view(-1, 28, 28)) # reshape x to (batch, time_step, input_size)
        b_y = Variable(y) # batch y

        output = rnn(b_x) # rnn output
        loss = loss_func(output, b_y) # cross entropy loss
        optimizer.zero_grad() # clear gradients for this training step
        loss.backward() # backpropagation, compute gradients
        optimizer.step() # apply gradients

    if step % 50 == 0:
        test_output = rnn(test_x) # (samples, time_step, input_size)
        pred_y = torch.max(test_output, 1)[1].data.numpy().squeeze()
        accuracy = sum(pred_y == test_y) / float(test_y.size)
        print('Epoch: ', epoch, '| train loss: %.4f' % loss.data, '| test accuracy: %.2f' % accuracy)

```

2) change the number of hidden nodes in the program three times and compare the results.

#3-2 hidden\_size=64

```

class RNN(nn.Module):
    def __init__(self):
        super(RNN, self).__init__()

        self.rnn = nn.LSTM( # if use nn.RNN(), it hardly learns
            input_size=INPUT_SIZE,
            hidden_size=64, # rnn hidden unit
            num_layers=1, # number of rnn layer
            batch_first=True, # input & output will has batch size as 1s dimension. e.g. (batch,
time_step, input_size)
        )

        self.out = nn.Linear(64, 10)

    def forward(self, x):
        # x shape (batch, time_step, input_size)
        # r_out shape (batch, time_step, output_size)
        # h_n shape (n_layers, batch, hidden_size)
        # h_c shape (n_layers, batch, hidden_size)
        r_out, (h_n, h_c) = self.rnn(x, None) # None represents zero initial hidden state

        # choose r_out at the last time step
        out = self.out(r_out[:, -1, :])
        return out

rnn = RNN()
print(rnn)

```

```
optimizer = torch.optim.Adam(rnn.parameters(), lr=LR) # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()                    # the target label is not one-hot
```

```
RNN(
  (rnn): LSTM(28, 64, batch_first=True)
  (out): Linear(in_features=64, out_features=10, bias=True)
)
```

Epoch: 0	train loss: 2.2883	test accuracy: 0.10
Epoch: 0	train loss: 0.8980	test accuracy: 0.56
Epoch: 0	train loss: 1.0743	test accuracy: 0.70
Epoch: 0	train loss: 0.6517	test accuracy: 0.83
Epoch: 0	train loss: 0.5668	test accuracy: 0.83
Epoch: 0	train loss: 0.3297	test accuracy: 0.88
Epoch: 0	train loss: 0.4544	test accuracy: 0.89
Epoch: 0	train loss: 0.3315	test accuracy: 0.92
Epoch: 0	train loss: 0.1421	test accuracy: 0.92
Epoch: 0	train loss: 0.3268	test accuracy: 0.93
Epoch: 0	train loss: 0.0576	test accuracy: 0.93
Epoch: 0	train loss: 0.2015	test accuracy: 0.94
Epoch: 0	train loss: 0.1035	test accuracy: 0.93
Epoch: 0	train loss: 0.1204	test accuracy: 0.94
Epoch: 0	train loss: 0.1826	test accuracy: 0.94
Epoch: 0	train loss: 0.1199	test accuracy: 0.95
Epoch: 0	train loss: 0.0627	test accuracy: 0.94
Epoch: 0	train loss: 0.1477	test accuracy: 0.96
Epoch: 0	train loss: 0.2381	test accuracy: 0.95

### #3-2 hidden\_size=32

```
class RNN(nn.Module):
    def __init__(self):
        super(RNN, self).__init__()

        self.rnn = nn.LSTM(          # if use nn.RNN(), it hardly learns
            input_size=INPUT_SIZE,
            hidden_size=32,          # rnn hidden unit
            num_layers=1,            # number of rnn layer
            batch_first=True,        # input & output will has batch size as 1s dimension. e.g. (batch,
            time_step, input_size)
        )

        self.out = nn.Linear(32, 10)
```

```
def forward(self, x):
    # x shape (batch, time_step, input_size)
    # r_out shape (batch, time_step, output_size)
    # h_n shape (n_layers, batch, hidden_size)
    # h_c shape (n_layers, batch, hidden_size)
    r_out, (h_n, h_c) = self.rnn(x, None)    # None represents zero initial hidden state

    # choose r_out at the last time step
    out = self.out(r_out[:, -1, :])
    return out
```

```
rnn = RNN()
print(rnn)
```

```
optimizer = torch.optim.Adam(rnn.parameters(), lr=LR) # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()                    # the target label is not one-hot
```



```
RNN(
    (rnn): LSTM(28, 32, batch_first=True)
    (out): Linear(in_features=32, out_features=10, bias=True)
)
```

/Users/soojinlee/opt/anaconda3/lib/python3.7/site-packages/ip.  
d and now has no effect. Use `with torch.no\_grad():` instead.

```
Epoch: 0 | train loss: 2.3050 | test accuracy: 0.12
Epoch: 0 | train loss: 1.3808 | test accuracy: 0.49
Epoch: 0 | train loss: 1.0123 | test accuracy: 0.66
Epoch: 0 | train loss: 0.6900 | test accuracy: 0.77
Epoch: 0 | train loss: 0.7815 | test accuracy: 0.79
Epoch: 0 | train loss: 0.3825 | test accuracy: 0.81
Epoch: 0 | train loss: 0.5824 | test accuracy: 0.85
Epoch: 0 | train loss: 0.3229 | test accuracy: 0.89
Epoch: 0 | train loss: 0.1642 | test accuracy: 0.89
Epoch: 0 | train loss: 0.3344 | test accuracy: 0.90
Epoch: 0 | train loss: 0.4474 | test accuracy: 0.92
Epoch: 0 | train loss: 0.2098 | test accuracy: 0.93
Epoch: 0 | train loss: 0.5863 | test accuracy: 0.92
Epoch: 0 | train loss: 0.2316 | test accuracy: 0.92
Epoch: 0 | train loss: 0.0935 | test accuracy: 0.94
Epoch: 0 | train loss: 0.2247 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1834 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1711 | test accuracy: 0.94
Epoch: 0 | train loss: 0.3964 | test accuracy: 0.95
```

#3-2 hidden\_size=16

```
class RNN(nn.Module):
    def __init__(self):
        super(RNN, self).__init__()

        self.rnn = nn.LSTM(          # if use nn.RNN(), it hardly learns
            input_size=INPUT_SIZE,
            hidden_size=16,          # rnn hidden unit
            num_layers=1,            # number of rnn layer
            batch_first=True,        # input & output will has batch size as 1s dimension. e.g. (batch,
time_step, input_size)
        )

        self.out = nn.Linear(16, 10)

    def forward(self, x):
        # x shape (batch, time_step, input_size)
        # r_out shape (batch, time_step, output_size)
        # h_n shape (n_layers, batch, hidden_size)
        # h_c shape (n_layers, batch, hidden_size)
        r_out, (h_n, h_c) = self.rnn(x, None) # None represents zero initial hidden state

        # choose r_out at the last time step
        out = self.out(r_out[:, -1, :])
        return out

rnn = RNN()
print(rnn)

optimizer = torch.optim.Adam(rnn.parameters(), lr=LR) # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()                    # the target label is not one-hot
```

```
RNN(
    (rnn): LSTM(28, 16, batch_first=True)
    (out): Linear(in_features=16, out_features=10, bias=True)
)
```

/Users/soojinlee/opt/anaconda3/lib/python3.7/site-packages/ip  
d and now has no effect. Use `with torch.no\_grad():` instead.

```
Epoch: 0 | train loss: 2.3317 | test accuracy: 0.11
Epoch: 0 | train loss: 1.3542 | test accuracy: 0.45
Epoch: 0 | train loss: 1.0651 | test accuracy: 0.54
Epoch: 0 | train loss: 1.0748 | test accuracy: 0.60
Epoch: 0 | train loss: 0.6775 | test accuracy: 0.70
Epoch: 0 | train loss: 0.7007 | test accuracy: 0.74
Epoch: 0 | train loss: 0.6877 | test accuracy: 0.71
Epoch: 0 | train loss: 0.6917 | test accuracy: 0.79
Epoch: 0 | train loss: 0.3936 | test accuracy: 0.82
Epoch: 0 | train loss: 0.3910 | test accuracy: 0.83
Epoch: 0 | train loss: 0.4074 | test accuracy: 0.85
Epoch: 0 | train loss: 0.3370 | test accuracy: 0.86
Epoch: 0 | train loss: 0.3332 | test accuracy: 0.86
Epoch: 0 | train loss: 0.3917 | test accuracy: 0.88
Epoch: 0 | train loss: 0.3024 | test accuracy: 0.88
Epoch: 0 | train loss: 0.4465 | test accuracy: 0.89
Epoch: 0 | train loss: 0.3496 | test accuracy: 0.88
Epoch: 0 | train loss: 0.3620 | test accuracy: 0.87
Epoch: 0 | train loss: 0.2999 | test accuracy: 0.88
```

## Report

⇒ The smaller the hidden size, the lower the accuracy.

3) change the current optimization method to other optimization methods

#3-3 Adam

```
class RNN(nn.Module):
    def __init__(self):
        super(RNN, self).__init__()

        self.rnn = nn.LSTM(          # if use nn.RNN(), it hardly learns
            input_size=INPUT_SIZE,
            hidden_size=64,           # rnn hidden unit
            num_layers=1,             # number of rnn layer
            batch_first=True,         # input & output will has batch size as 1s dimension. e.g. (batch,
time_step, input_size)
        )

        self.out = nn.Linear(64, 10)

    def forward(self, x):
        # x shape (batch, time_step, input_size)
        # r_out shape (batch, time_step, output_size)
        # h_n shape (n_layers, batch, hidden_size)
        # h_c shape (n_layers, batch, hidden_size)
        r_out, (h_n, h_c) = self.rnn(x, None) # None represents zero initial hidden state

        # choose r_out at the last time step
        out = self.out(r_out[:, -1, :])
        return out
```

```

rnn = RNN()
print(rnn)

optimizer = torch.optim.Adam(rnn.parameters(), lr=LR)    # optimize all rnn parameters
loss_func = nn.CrossEntropyLoss()                        # the target label is not one-hot

```

```

RNN(
  (rnn): LSTM(28, 64, batch_first=True)
  (out): Linear(in_features=64, out_features=10, bias=True)
)

```

```

Epoch: 0 | train loss: 2.2883 | test accuracy: 0.10
Epoch: 0 | train loss: 0.8980 | test accuracy: 0.56
Epoch: 0 | train loss: 1.0743 | test accuracy: 0.70
Epoch: 0 | train loss: 0.6517 | test accuracy: 0.83
Epoch: 0 | train loss: 0.5668 | test accuracy: 0.83
Epoch: 0 | train loss: 0.3297 | test accuracy: 0.88
Epoch: 0 | train loss: 0.4544 | test accuracy: 0.89
Epoch: 0 | train loss: 0.3315 | test accuracy: 0.92
Epoch: 0 | train loss: 0.1421 | test accuracy: 0.92
Epoch: 0 | train loss: 0.3268 | test accuracy: 0.93
Epoch: 0 | train loss: 0.0576 | test accuracy: 0.93
Epoch: 0 | train loss: 0.2015 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1035 | test accuracy: 0.93
Epoch: 0 | train loss: 0.1204 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1826 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1199 | test accuracy: 0.95
Epoch: 0 | train loss: 0.0627 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1477 | test accuracy: 0.96
Epoch: 0 | train loss: 0.2381 | test accuracy: 0.95

```

### #3-3 optimizer – RMSprop

```

class RNN(nn.Module):
    def __init__(self):
        super(RNN, self).__init__()

        self.rnn = nn.LSTM(          # if use nn.RNN(), it hardly learns
            input_size=INPUT_SIZE,
            hidden_size=64,           # rnn hidden unit
            num_layers=1,             # number of rnn layer
            batch_first=True,         # input & output will has batch size as 1s dimension. e.g. (batch,
time_step, input_size)
        )

        self.out = nn.Linear(64, 10)

    def forward(self, x):
        # x shape (batch, time_step, input_size)
        # r_out shape (batch, time_step, output_size)
        # h_n shape (n_layers, batch, hidden_size)
        # h_c shape (n_layers, batch, hidden_size)
        r_out, (h_n, h_c) = self.rnn(x, None)    # None represents zero initial hidden state

        # choose r_out at the last time step
        out = self.out(r_out[:, -1, :])
        return out

rnn = RNN()
print(rnn)

```

```
optimizer = torch.optim.RMSprop(rnn.parameters(), lr=LR) # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss() # the target label is not one-hot
```

```
RNN(
    (rnn): LSTM(28, 64, batch_first=True)
    (out): Linear(in_features=64, out_features=10, bias=True)
)
```

/Users/soojinlee/opt/anaconda3/lib/python3.7/site-packages/ipython and now has no effect. Use `with torch.no\_grad():` instead.

```
Epoch: 0 | train loss: 2.2883 | test accuracy: 0.10
Epoch: 0 | train loss: 1.2381 | test accuracy: 0.43
Epoch: 0 | train loss: 0.8018 | test accuracy: 0.64
Epoch: 0 | train loss: 0.7949 | test accuracy: 0.74
Epoch: 0 | train loss: 0.4477 | test accuracy: 0.78
Epoch: 0 | train loss: 0.3670 | test accuracy: 0.84
Epoch: 0 | train loss: 0.4815 | test accuracy: 0.87
Epoch: 0 | train loss: 0.5720 | test accuracy: 0.85
Epoch: 0 | train loss: 0.3554 | test accuracy: 0.89
Epoch: 0 | train loss: 0.3389 | test accuracy: 0.90
Epoch: 0 | train loss: 0.1689 | test accuracy: 0.92
Epoch: 0 | train loss: 0.2035 | test accuracy: 0.91
Epoch: 0 | train loss: 0.1487 | test accuracy: 0.93
Epoch: 0 | train loss: 0.2166 | test accuracy: 0.92
Epoch: 0 | train loss: 0.2518 | test accuracy: 0.93
Epoch: 0 | train loss: 0.1977 | test accuracy: 0.92
Epoch: 0 | train loss: 0.2616 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1658 | test accuracy: 0.94
Epoch: 0 | train loss: 0.3873 | test accuracy: 0.94
```

### #3-3 optimizer – Adadelata

```
class RNN(nn.Module):
    def __init__(self):
        super(RNN, self).__init__()

        self.rnn = nn.LSTM(          # if use nn.RNN(), it hardly learns
            input_size=INPUT_SIZE,
            hidden_size=64,          # rnn hidden unit
            num_layers=1,            # number of rnn layer
            batch_first=True,        # input & output will has batch size as 1s dimension. e.g. (batch,
time_step, input_size)
        )

        self.out = nn.Linear(64, 10)

    def forward(self, x):
        # x shape (batch, time_step, input_size)
        # r_out shape (batch, time_step, output_size)
        # h_n shape (n_layers, batch, hidden_size)
        # h_c shape (n_layers, batch, hidden_size)
        r_out, (h_n, h_c) = self.rnn(x, None) # None represents zero initial hidden state

        # choose r_out at the last time step
        out = self.out(r_out[:, -1, :])
        return out
```

```

rnn = RNN()
print(rnn)

optimizer = torch.optim.Adadelta(rnn.parameters(), lr=LR)    # optimize all rnn parameters
loss_func = nn.CrossEntropyLoss()                            # the target label is not one-hot

```

```

RNN(
  (rnn): LSTM(28, 64, batch_first=True)
  (out): Linear(in_features=64, out_features=10, bias=True)
)

```

/Users/soojinlee/opt/anaconda3/lib/python3.7/site-packages/ipynb and now has no effect. Use `with torch.no\_grad():` instead.

```

Epoch: 0 | train loss: 2.2883 | test accuracy: 0.08
Epoch: 0 | train loss: 2.3065 | test accuracy: 0.08
Epoch: 0 | train loss: 2.3140 | test accuracy: 0.09
Epoch: 0 | train loss: 2.3099 | test accuracy: 0.09
Epoch: 0 | train loss: 2.3207 | test accuracy: 0.09
Epoch: 0 | train loss: 2.2916 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3055 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3028 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3026 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3081 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3226 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3133 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3073 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3168 | test accuracy: 0.10
Epoch: 0 | train loss: 2.2932 | test accuracy: 0.10
Epoch: 0 | train loss: 2.2968 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3007 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3038 | test accuracy: 0.10
Epoch: 0 | train loss: 2.3010 | test accuracy: 0.10

```

## Report

⇒ adadelta cannot be used because it is very inaccurate to use. rmsprop is also good accuracy, but adam's accuracy is the best.

4) change LSTM to GRU (or vice versa). Compare the results

```

class RNN(nn.Module):
    def __init__(self):
        super(RNN, self).__init__()

        self.rnn = nn.GRU(          # if use nn.RNN(), it hardly learns
            input_size=INPUT_SIZE,
            hidden_size=64,          # rnn hidden unit
            num_layers=2,            # number of rnn layer
            batch_first=True,        # input & output will has batch size as 1s dimension. e.g. (batch,
time_step, input_size)
        )

        self.out = nn.Linear(64, 10)

    def forward(self, x):
        # x shape (batch, time_step, input_size)
        # r_out shape (batch, time_step, output_size)
        # h_n shape (n_layers, batch, hidden_size)
        # h_c shape (n_layers, batch, hidden_size)
        r_out, (h_n, h_c) = self.rnn(x, None)    # None represents zero initial hidden state

```

```

        # choose r_out at the last time step
        out = self.out(r_out[:, -1, :])
        return out

rnn = RNN()
print(rnn)

optimizer = torch.optim.Adam(rnn.parameters(), lr=LR)    # optimize all rnn parameters
loss_func = nn.CrossEntropyLoss()                        # the target label is not one-hot

```

```

RNN(
  (rnn): GRU(28, 64, num_layers=2, batch_first=True)
  (out): Linear(in_features=64, out_features=10, bias=True)
)

```

```

/Users/soojinlee/opt/anaconda3/lib/python3.7/site-packages/i
d and now has no effect. Use `with torch.no_grad():` instead

```

Epoch: 0	train loss: 2.3132	test accuracy: 0.12
Epoch: 0	train loss: 0.8244	test accuracy: 0.67
Epoch: 0	train loss: 0.3006	test accuracy: 0.85
Epoch: 0	train loss: 0.2121	test accuracy: 0.90
Epoch: 0	train loss: 0.1530	test accuracy: 0.92
Epoch: 0	train loss: 0.3483	test accuracy: 0.92
Epoch: 0	train loss: 0.0755	test accuracy: 0.94
Epoch: 0	train loss: 0.2027	test accuracy: 0.95
Epoch: 0	train loss: 0.3304	test accuracy: 0.95
Epoch: 0	train loss: 0.0169	test accuracy: 0.96
Epoch: 0	train loss: 0.2831	test accuracy: 0.93
Epoch: 0	train loss: 0.1494	test accuracy: 0.96
Epoch: 0	train loss: 0.2699	test accuracy: 0.96
Epoch: 0	train loss: 0.0510	test accuracy: 0.96
Epoch: 0	train loss: 0.2100	test accuracy: 0.93
Epoch: 0	train loss: 0.1083	test accuracy: 0.95
Epoch: 0	train loss: 0.0791	test accuracy: 0.97
Epoch: 0	train loss: 0.0573	test accuracy: 0.97
Epoch: 0	train loss: 0.3135	test accuracy: 0.96

## Report

⇒ GRU's accuracy is better than LSTM

5) choose ONE other parameters of RNN program (e.g. batch\_size, epochs, etc). Change the value of this parameter and compare the results.

#3-5 EPOCH=2

```

import torch
from torch import nn
from torch.autograd import Variable
import torchvision.datasets as dsets
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

```

```

torch.manual_seed(1)    # reproducible

```

# Hyper Parameters

```

EPOCH = 2    # train the training data n times, to save time, we just train 1 epoch

```

```

BATCH_SIZE = 64

```

```

TIME_STEP = 28    # rnn time step / image height

```

```

INPUT_SIZE = 28          # rnn input size / image width
LR = 0.01                # learning rate
DOWNLOAD_MNIST = False   # set to True if haven't download the data

# Mnist digital dataset
train_data = datasets.MNIST(
    root='./mnist/',
    train=True,          # this is training data
    transform=transforms.ToTensor(), # Converts a PIL.Image or numpy.ndarray to
                                   # torch.FloatTensor of shape (C x H x W) and normalize in the
range [0.0, 1.0]
    download=DOWNLOAD_MNIST, # download it if you don't have it
)

# plot one example
print(train_data.train_data.size()) # (60000, 28, 28)
print(train_data.train_labels.size()) # (60000)
plt.imshow(train_data.train_data[0].numpy(), cmap='gray')
plt.title('%i' % train_data.train_labels[0])
plt.show()

# Data Loader for easy mini-batch return in training
train_loader = torch.utils.data.DataLoader(dataset=train_data, batch_size=BATCH_SIZE, shuffle=True)

# convert test data into Variable, pick 2000 samples to speed up testing
test_data = datasets.MNIST(root='./mnist/', train=False, transform=transforms.ToTensor())
test_x = Variable(test_data.test_data, volatile=True).type(torch.FloatTensor)[:2000]/255. # shape (2000, 28,
28) value in range(0,1)
test_y = test_data.test_labels.numpy().squeeze()[:2000] # covert to numpy array

class RNN(nn.Module):
    def __init__(self):
        super(RNN, self).__init__()

        self.rnn = nn.LSTM( # if use nn.RNN(), it hardly learns
            input_size=INPUT_SIZE,
            hidden_size=64, # rnn hidden unit
            num_layers=1, # number of rnn layer
            batch_first=True, # input & output will has batch size as 1s dimension. e.g. (batch,
time_step, input_size)
        )

        self.out = nn.Linear(64, 10)

    def forward(self, x):
        # x shape (batch, time_step, input_size)
        # r_out shape (batch, time_step, output_size)
        # h_n shape (n_layers, batch, hidden_size)
        # h_c shape (n_layers, batch, hidden_size)
        r_out, (h_n, h_c) = self.rnn(x, None) # None represents zero initial hidden state

        # choose r_out at the last time step
        out = self.out(r_out[:, -1, :])
        return out

rnn = RNN()
print(rnn)

```

```

optimizer = torch.optim.Adam(rnn.parameters(), lr=LR)    # optimize all cnn parameters
loss_func = nn.CrossEntropyLoss()                      # the target label is not one-hot

# training and testing
for epoch in range(EPOCH):
    for step, (x, y) in enumerate(train_loader):         # gives batch data
        b_x = Variable(x.view(-1, 28, 28))             # reshape x to (batch, time_step, input_size)
        b_y = Variable(y)                              # batch y

        output = rnn(b_x)                              # rnn output
        loss = loss_func(output, b_y)                  # cross entropy loss
        optimizer.zero_grad()                          # clear gradients for this training step
        loss.backward()                                # backpropagation, compute gradients
        optimizer.step()                               # apply gradients

```

```

if step % 50 == 0:
    test_output = rnn(test_x)                          # (samples, time_step, input_size)
    pred_y = torch.max(test_output, 1)[1].data.numpy().squeeze()
    accuracy = sum(pred_y == test_y) / float(test_y.size)
    print('Epoch: ', epoch, '| train loss: %.4f' % loss.data, '| test accuracy: %.2f' % accuracy)

```

```

Epoch: 0 | train loss: 0.2015 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1035 | test accuracy: 0.93
Epoch: 0 | train loss: 0.1204 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1826 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1199 | test accuracy: 0.95
Epoch: 0 | train loss: 0.0627 | test accuracy: 0.94
Epoch: 0 | train loss: 0.1477 | test accuracy: 0.96
Epoch: 0 | train loss: 0.2381 | test accuracy: 0.95
Epoch: 1 | train loss: 0.0900 | test accuracy: 0.95
Epoch: 1 | train loss: 0.0148 | test accuracy: 0.95
Epoch: 1 | train loss: 0.0904 | test accuracy: 0.95
Epoch: 1 | train loss: 0.1294 | test accuracy: 0.95
Epoch: 1 | train loss: 0.1356 | test accuracy: 0.96
Epoch: 1 | train loss: 0.1203 | test accuracy: 0.96
Epoch: 1 | train loss: 0.0942 | test accuracy: 0.96
Epoch: 1 | train loss: 0.1747 | test accuracy: 0.96
Epoch: 1 | train loss: 0.0608 | test accuracy: 0.96
Epoch: 1 | train loss: 0.1119 | test accuracy: 0.96
Epoch: 1 | train loss: 0.0941 | test accuracy: 0.96
Epoch: 1 | train loss: 0.0280 | test accuracy: 0.96
Epoch: 1 | train loss: 0.1034 | test accuracy: 0.96
Epoch: 1 | train loss: 0.2348 | test accuracy: 0.97
Epoch: 1 | train loss: 0.3485 | test accuracy: 0.96
Epoch: 1 | train loss: 0.1760 | test accuracy: 0.97
Epoch: 1 | train loss: 0.0346 | test accuracy: 0.97
Epoch: 1 | train loss: 0.0973 | test accuracy: 0.97
Epoch: 1 | train loss: 0.1013 | test accuracy: 0.97

```

메도

## Report

⇒ By increase Epoch as 2, accuracy was increased. But running time is also increased.

6) compare the accuracy of RNN for Mnist with that of CNN.

CNN's accuracy is higher than RNN on average. Also CNN's train loss is lower than RNN. CNN shows better function.